

Workshop 8 TEMPLATE

12pm

Michael Torbensen

2026-05-22

Table of contents

1. Set up	2
2. Clean data	3
3. Visualizing data	5
What is the relationship between SLA and plant mass?	5
What are the differences in total mass between water treatments?	6
What are the differences in total mass between species?	7
4. Fitting models	9
Model fitting	10
Model diagnostics	11
Model selection	12
Stop and think: what does this model mean?	14
4. Model predictions	14
5. Final figure	15
6. Writing	16
Figure caption	16
Basic components of test	16
Written communication	17
Methods	17
Results	17

In this workshop, we will answer the question: How do specific leaf area, water treatment, and species influence plant mass?

Specific leaf area is a continuous variable measured in cm^2/g .

Water treatment is a categorical variable (i.e. a *factor*) with 2 levels: drought stressed (DS) and well watered (WW).

Species is a categorical variable (again, a *factor*) with 6 levels.

Species name	Species code	Common name
<i>Encelia californica</i>	ENCCAL	Bush sunflower
<i>Eschsholzia californica</i>	ESCCAL	California poppy
<i>Penstemon centranthifolius</i>	PENCEN	Scarlet bugler
<i>Grindelia camporum</i>	GRINCAM	Gumweed
<i>Salvia leucophylla</i>	SALLEU	Purple sage
<i>Stipa pulchra</i>	STIPUL	Purple needlegrass
<i>Lotus scoparius</i>	LOTSCO	Deerweed

1. Set up

```
# read in your packages here
library(tidyverse) # general use
library(here) # file organization
library(janitor) # cleaning data frames
library(readxl) # reading excel files
library(scales) # modifying axis labels
library(ggeffects) # getting model predictions
library(MuMin) # model selection

# read in your data here

drought_exp <- read_xlsx(
  here("data", "Valliere_etal_EcoApps_Data.xlsx"),
  sheet = "First Harvest"
)
```

```
# storing colors to use for species
lotsco_col <- "#E69512"
pencen_col <- "#D6264F"
salleu_col <- "#6D397D"
enccal_col <- "#3A5565"
stipul_col <- "#3F564F"
```

```

escscal_col <- "#515481"
gricam_col <- "#6C91BD"

# storing colors to use for water treatments
ds_col <- "#A62F03"
ww_col <- "#045CB4"

# storing a ggplot theme (that will be used for all ggplots)
theme_set(theme_bw())

```

2. Clean data

Data source: Valliere, Justin; Zhang, Jacqueline; Sharifi, M.; Rundel, Philip (2019). Data from: Can we condition native plants to increase drought tolerance and improve restoration success? [Dataset]. Dryad. <https://doi.org/10.5061/dryad.v0861f7>

```

# creating new object
drought_exp_clean <- drought_exp |>
  # lowercase and removing spaces/adding only underscores
  clean_names() |>
  # creating species name
  mutate(species_name = recode_values(
    species,
    "ENCCAL" ~ "Encelia californica", # bush sunflower
    "ESCCAL" ~ "Eschscholzia californica", # California poppy
    "PENCEN" ~ "Penstemon centranthifolius", # Scarlet bugler
    "GRICAM" ~ "Grindelia camporum", # Gumweed
    "SALLEU" ~ "Salvia leucophylla", # purple sage
    "STIPUL" ~ "Stipa pulchra", # purple needlegrass
    "LOTSCO" ~ "Lotus scoparius" # deerweed
  )) |>
  # giving full name to water type
  mutate(water_treatment = recode_values(
    water,
    "WW" ~ "Well watered",
    "DS" ~ "Drought stressed"
  )) |>
  # changing the order of each species, organizing the names and data within columns
  mutate(species_name = as_factor(species_name),
    # factors that show first are always going to be the reference level
    species_name = fct_relevel(species_name,

```

```

        "Lotus scoparius",
        "Penstemon centranthifolius",
        "Salvia leucophylla",
        "Encelia californica",
        "Stipa pulchra",
        "Eschscholzia californica",
        "Grindelia camporum")) |>
# reordering water_treatment levels
mutate(water_treatment = as_factor(water_treatment),
       water_treatment = fct_relevel(water_treatment,
                                     "Drought stressed",
                                     "Well watered")) |>
# selecting columns of interest
select(species_name, water_treatment, sla, total_g)

```

Insert code to double check the structure of the data frame here.

```
str(drought_exp_clean)
```

```

tibble [70 x 4] (S3: tbl_df/tbl/data.frame)
 $ species_name   : Factor w/ 7 levels "Lotus scoparius",...: 4 4 4 4 4 4 4 4 4 4 ...
 $ water_treatment: Factor w/ 2 levels "Drought stressed",...: 2 2 2 2 2 1 1 1 1 1 ...
 $ sla            : num [1:70] 170 215 209 216 222 ...
 $ total_g        : num [1:70] 0.455 0.329 0.45 0.359 0.352 ...

```

Insert code to display 10 random rows from the data frame here.

```

slice_sample(
  drought_exp_clean, # clean data frame to display 10 rows
  n = 10
)

```

```

# A tibble: 10 x 4
  species_name      water_treatment    sla total_g
  <fct>            <fct>          <dbl>  <dbl>
1 Stipa pulchra    Drought stressed  146.   0.246
2 Encelia californica Drought stressed  161.   0.273
3 Encelia californica Drought stressed  192.   0.261
4 Salvia leucophylla Drought stressed   77.6  0.186
5 Eschscholzia californica Well watered      117.   0.574

```

6	Penstemon centranthifolius	Drought stressed	146.	0.126
7	Salvia leucophylla	Drought stressed	74.4	0.131
8	Lotus scoparius	Drought stressed	116.	0.0802
9	Salvia leucophylla	Drought stressed	96.1	0.278
10	Eschscholzia californica	Well watered	116.	0.423

3. Visualizing data


What is the relationship between SLA and plant mass?

Make a scatterplot to visualize SLA (the predictor) on the x-axis, and total mass (the response) on the y-axis.

Label the axes with units where appropriate.

Add a title describing the relationship (e.g. as SLA increases, total mass _____).

```
# base layer: ggplot
ggplot(data = drought_exp_clean,
       # x-axis: specific leaf area
       # y-axis: total mass of leaf
       mapping = aes(x = sla,
                     y = total_g)) +
  # adding points to plot
  geom_point() +
  # labeling x- and y-axis and title
  labs(x = "Specific leaf Area (cm2 / g)",
       y = "total leaf mass in g",
       title = "As SLA Increases, Total g Increases")
```



What are the differences in total mass between water treatments?

Make a jitterplot to visualize the differences in plant mass between water treatments.

Label the axes with units where appropriate.


Add a title describing the differences in (mean or median) plant mass between water treatments.

```
# base laser: ggplot
ggplot(data = drought_exp_clean,
  # x-axis: treatment
  # y-axis = treatment type
  # coloring by treatment
  mapping = aes(x = water_treatment,
    y = total_g,
    color = water_treatment)) +
# representing individual points
geom_jitter(height = 0,
  width = 0.1) +
scale_color_manual(values = c(
  "Drought stressed" = ds_col,
  "Well Watered" = ww_col
```

```

)) +
labs(x = "water treatment type",
     y = "total leaf mass in g",
     title = "Plants well watered have higher mass than drought-stressed")

```



What are the differences in total mass between species?

Make a jitterplot to visualize the differences in plant mass between species.

Label the axes with units where appropriate.

Add a title describing the differences in (mean or median) plant mass between species.


Bonus: order the x-axis by decreasing mean plant mass.

```

ggplot(data = drought_exp_clean,
       mapping = aes(x = species_name,
                     y = total_g,
                     color = species_name)) + # coloring by species
  # represent individuals in each species
  geom_jitter(height = 0) +
  # relabeling x- and y-axis and title
  labs(x = "Species Name",
       y = "Total leaf mass in g",
       title = "Species differ in plant mass") +
  # wrapping x-axis text
  scale_x_discrete(labels = label_wrap(10)) +
  # using custom colors
  scale_color_manual(values = c(
    "Lotus scoparius" = lotsco_col,
    "Pentstemon centranthifolius" = pencen_col,
    "Salvia leucophylla" = sallev_col,
    "Encelia californica" = enccal_col,
    "Stipa pulchra" = stipul_col,
    "Eschscholzia californica" = esccal_col,
    "Grindelia camporum" = gricam_col)) +
  # taking out legend
  theme(legend.position = "none")

```

Species differ in plant mass



4. Fitting models

8 models total:

Model number	SLA	Water treatment	Species	Predictor list
0				no predictors (null model)
1	X	X	X	all predictors (full model)
2	X	X		SLA and water treatment
3	X		X	SLA and species
4		X	X	water treatment and species
5	X			SLA
6		X		water treatment
7			X	species

Model fitting

```
# model 0: null model

model0 <- lm(
  total_g ~ 1,
  data = drought_exp_clean
)

# model 1: all predictors

model1 <- lm(
  total_g ~ sla + water_treatment + species_name,
  data = drought_exp_clean
)

# model 2: SLA and water treatment

model2 <- lm(
  total_g ~ sla + water_treatment,
  data = drought_exp_clean
)

# model 3: SLA and species

model3 <- lm(
  total_g ~ sla + species_name,
  data = drought_exp_clean
)

# model 4: water treatment and species

model4 <- lm(
  total_g ~ water_treatment + species_name,
  data = drought_exp_clean
)


# model 5: SLA

model5 <- lm(
  total_g ~ sla,
  data = drought_exp_clean
)
```

```
)  
  
# model 6: water treatment  
  
model6 <- lm(  
  total_g ~ water_treatment,  
  data = drought_exp_clean  
)  
  
# model 7: species  
  
model7 <- lm(  
  total_g ~ species_name,  
  data = drought_exp_clean  
)
```

Model diagnostics

```
# displaying plots in 2x2 grid  
par(mfrow = c(2, 2))  
# plotting diagnostics  
plot(model4)
```



- residuals are normally distributed from QQ plot
- residuals are homoscedastic from residuals vs fitted and scale-location plots
- no outliers that could be influencing model estimates (from our leverage plot)

Model selection

Insert code to use the `AICc()` function to calculate the AIC for each model.

Arrange by decreasing AIC.

```
AICc(
  model0,
  model1,
  model2,
  model3,
  model4,
  model5,
  model6,
  model7
) |>
  arrange(AICc)
```

	df	AICc
model4	9	-156.19595
model11	10	-153.75361
model3	9	-124.07569
model7	8	-120.30191
model2	4	-95.82521
model5	3	-88.85180
model6	3	-86.76661
model0	2	-74.98036

Insert code to display the summary for the best model.

```
summary(
  model4
)
```

Call:

```
lm(formula = total_g ~ water_treatment + species_name, data = drought_exp_clean)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.157087	-0.046953	-0.003733	0.041244	0.192657

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.05455	0.02451	2.225	0.02973 *
water_treatmentWell watered	0.11695	0.01733	6.746	5.90e-
09 ***				
species_namePenstemon centranthifolius	0.05003	0.03243	1.543	0.12799
species_nameSalvia leucophylla	0.12020	0.03243	3.706	0.00045 ***
species_nameEncelia californica	0.21774	0.03243	6.714	6.70e-
09 ***				
species_nameStipa pulchra	0.22881	0.03243	7.055	1.72e-
09 ***				
species_nameEschscholzia californica	0.23164	0.03243	7.143	1.22e-
09 ***				
species_nameGrindelia camporum	0.31335	0.03243	9.662	5.53e-
14 ***				

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.07252 on 62 degrees of freedom
Multiple R-squared: 0.7535, Adjusted R-squared: 0.7257
F-statistic: 27.08 on 7 and 62 DF, p-value: < 2.2e-16

Stop and think: what does this model mean?

What is the best model?

The best model includes

How much variation in the response (total mass, in grams) does this model explain?

[insert your response here]

How do we interpret the effects of the predictors on the response variable (again, total mass in grams)?

- Species influences plant mass.
- Treatment influences plant mass.
- (Intercept): the intercept is the model prediction for plant mass for *Lotus scoparius* in the drought stressed treatment.
- `water_treatmentWell watered`: model prediction for the difference in *Lotus scoparius* mass between well watered and drought stressed treatment.
- `species_namepenstemon centranthifolius`: difference between *penstemon* mass in the well watered treatment and *lotus scoparius* mass in well watered treatment
- `species_nameEschscholzia californica`: difference between *Eschscholzia californica* mass in well watered treatment and *lotus scoparius* mass in well watered treatment

4. Model predictions

Insert code to generate model predictions for the best model.

Rename the `x` column and `group` column to match the original data frame (`drought_exp_clean`).

Look at the object before plotting.

```

preds <- ggpredict(
  model4, # model object
  terms = c("water_treatment", "species_name") # predictors
) |>
  rename(water_treatment = x,
         species_name = group)

```


5. Final figure

Create a clean, publication ready figure depicting predictions for the best model and the underlying data.

```

# base layer: ggplot
ggplot(data = drought_exp_clean,
       mapping = aes(x = water_treatment,
                     y = total_g,
                     color = water_treatment)) +
# jitter: individual plats for each treatment in each species
geom_jitter(height = 0,
            width = 0.1,
            alpha = 0.4,
            shape = 21) +
# point range: model predictions with 95% CI
geom_pointrange(data = preds,
               mapping = aes(x = water_treatment,
                             y = predicted,
                             ymin = conf.low,
                             ymax = conf.high)) +
# creates different panels for each species
facet_wrap(~ species_name,
          nrow = 1) + # all plots in a single row
# adds custom colors
scale_color_manual(values = c(
  "Drought stressed" = ds_col,
  "Well watered" = ww_col
)) +
# removes legend
theme(legend.position = "none") +
# relabels x- and y-axis
labs(x = "Treatment",
     y = "Total plant mass in g")

```



6. Writing

Figure caption

Basic components of test

predictors

response

test

distribution

parameters (degrees of freedom)

R^2

p-value

significance level

reference level(s)

Written communication

Methods

We used linear models to determine the predictive relationship between specific leaf area (cm^2/g), treatment, and species on plant mass (g). We evaluated model assumptions of normally distributed errors using QQ plots of model residuals. We evaluated assumptions of constant variance of errors using Residuals vs Fitted and Scale-Location plots. We used Akaike's Information Criterion (AIC) to select the simplest model that explained the most variation in plant mass.

Results

We found that the simplest model describing the greatest variation in plant mass included treatment and species as predictors (linear model: $F(2, 67) = 14.57$, $p < 0.001$, $\alpha = 0.05$).